

2. German Credit Data

Hair Parra

2024-01-15

German Credit Data

- The German credit data set (Lichman, 2013) is another data set that we will use a lot.
- The data and description can be found here: [UCI Machine Learning Repository](#)
- This data set classifies 1000 people described by a set of 20 attributes as good or bad credit risks.
- The target variable, V21, is binary and is recorded to 0-1 (1-2 in the original data); 0=good risk and 1=bad risk.
- We work with a first version of the data set that includes 9 numeric covariates, 11 factor covariates, and the target variable.
- The goal of this example is to show you how to apply a **logistic regression Lasso model**

Data Preprocessing

```
# Prepare the German Credit data set
```

```
# Load path library  
library("here")
```

```
## here() starts at C:/Users/jairp/OneDrive/Desktop_remote/HEC Montreal/3. Winter 2024/Advanced Statistical Learning
```

```
# Set the path to the data file  
gercred = read.table(here("code_data_W2024", "german.data"))  
  
# Recode the target and two binary covariates to 0-1  
gercred$V21 = as.numeric(gercred$V21 == 2)  
gercred$V19 = as.numeric(gercred$V19 == "A192")  
gercred$V20 = as.numeric(gercred$V20 == "A201")  
  
# Convert factor variables to proper factors  
factor_vars = c("V1", "V3", "V4", "V6", "V7", "V9", "V10", "V12", "V14", "V15", "V17", "V18")  
gercred[factor_vars] <- lapply(gercred[factor_vars], factor)  
  
# Get the names of the factor variables  
fac_vars = vapply(gercred, is.factor, logical(1))  
namfac = names(fac_vars)[fac_vars]  
  
# Names of the numeric variables  
num_vars = vapply(gercred, is.numeric, logical(1))  
namnum = names(num_vars)[num_vars]  
  
# Display a summary of the German Credit data set  
summary(gercred)
```

```
##      V1          V2          V3          V4          V5          V6  
## A11:274   Min.    : 4.0   A30: 40   A43      :280   Min.    : 250   A61:603
```

```
## A12:269 1st Qu.:12.0 A31: 49 A40 :234 1st Qu.: 1366 A62:103
## A13: 63 Median :18.0 A32:530 A42 :181 Median : 2320 A63: 63
## A14:394 Mean :20.9 A33: 88 A41 :103 Mean : 3271 A64: 48
## 3rd Qu.:24.0 A34:293 A49 : 97 3rd Qu.: 3972 A65:183
## Max. :72.0 A46 : 50 Max. :18424
## (Other): 55
## V7 V8 V9 V10 V11 V12
## A71: 62 Min. :1.000 A91: 50 A101:907 Min. :1.000 A121:282
## A72:172 1st Qu.:2.000 A92:310 A102: 41 1st Qu.:2.000 A122:232
## A73:339 Median :3.000 A93:548 A103: 52 Median :3.000 A123:332
## A74:174 Mean :2.973 A94: 92 Mean :2.845 A124:154
## A75:253 3rd Qu.:4.000 3rd Qu.:4.000
## Max. :4.000 Max. :4.000
##
## V13 V14 V15 V16 V17 V18
## Min. :19.00 A141:139 A151:179 Min. :1.000 A171: 22 1:845
## 1st Qu.:27.00 A142: 47 A152:713 1st Qu.:1.000 A172:200 2:155
## Median :33.00 A143:814 A153:108 Median :1.000 A173:630
## Mean :35.55 Mean :1.407 A174:148
## 3rd Qu.:42.00 3rd Qu.:2.000
## Max. :75.00 Max. :4.000
##
## V19 V20 V21
## Min. :0.000 Min. :0.000 Min. :0.0
## 1st Qu.:0.000 1st Qu.:1.000 1st Qu.:0.0
## Median :0.000 Median :1.000 Median :0.0
## Mean :0.404 Mean :0.963 Mean :0.3
## 3rd Qu.:1.000 3rd Qu.:1.000 3rd Qu.:1.0
## Max. :1.000 Max. :1.000 Max. :1.0
##
```

Version with dummies

```
# load required libraries
library(fastDummies)
```

```
## Thank you for using fastDummies!
```

```
## To acknowledge our work, please cite the package:
```

```
## Kaplan, J. & Schlegel, B. (2023). fastDummies: Fast Creation of Dummy (Binary) Columns and Rows from Categorical Variables
```

```
# Create dummy variables for the factors
gercreddum=dummy_cols(gercred, remove_first_dummy=TRUE, remove_selected_columns=TRUE)

# Now all variables are numeric.
# There are 48 covariates and 1 binary target "V21".
# summary(gercreddum)
```

Train-test split

For the example, we create a training data set of size 600 and a test set of new data of size 400.

```
# Splitting the data into a training (ntrain=600) and a test (ntest=400) set

# Set the seed for reproducibility
```

```

set.seed(364565)

# Define the number of training and test samples
ntrain = 600
ntest = nrow(gercred) - ntrain

# Randomly select indices for the training set without replacement
indtrain = sample(1:nrow(gercred), ntrain, replace = FALSE)

# Create dummy variables for gercred data without the target variable (V21)
xdum = gercreddum # rename
xdum$V21 = NULL # target variable
xdum = as.matrix(xdum) # convert to matrix format

# Split the gercred data and dummy variables into training and test sets
gercredtrain = gercred[indtrain,]
gercredtest = gercred[-indtrain,]
gercreddumtrain = gercreddum[indtrain,]
gercreddumtest = gercreddum[-indtrain,]
gerxdumtrain = xdum[indtrain,]
gerxdumtest = xdum[-indtrain,]

```

Logistic Regression Lasso Model

Lasso:

$$\begin{aligned}\hat{\beta} &= \arg \min_{\beta} \left\{ \sum_{i=1}^n (y_i - (\beta_0 + \beta^T x_i))^2 + \lambda \sum_{j=1}^p |\beta_j| \right\} \\ &= \arg \min_{\beta} \|Y - X\beta\|_2^2 + \lambda \|\beta\|_1\end{aligned}$$

Elastic Net (likelihood-based)

$$\hat{\beta} = \arg \min_{\beta} \left\{ \frac{1}{n} \sum_{i=1}^n w_i \ell(y_i, (\beta_0 + \beta^T x_i))^2 + \lambda \left[(1 - \alpha) \frac{1}{2} \sum_{j=1}^p \beta_j^2 + \alpha \sum_{j=1}^p |\beta_j| \right] \right\}$$

Logistic Regression Elasticnet:

Using the equation above:

$$\hat{\beta} = \arg \min_{\beta} \left\{ \sum_{i=1}^n \left[y_i \log(1 + \exp(-\beta_0 - \beta^T x_i)) + (1 - y_i) \log \left(\frac{\exp(\beta_0 + \beta^T x_i)}{1 + \exp(\beta_0 + \beta^T x_i)} \right) \right] + \lambda \left[(1 - \alpha) \frac{1}{2} \sum_{j=1}^p \beta_j^2 + \alpha \sum_{j=1}^p |\beta_j| \right] \right\}$$


- $\alpha = 0$: **Ridge**
- $\alpha = 1$: **Lasso**

```
# Logistic regression with the lasso

# Set the seed for reproducibility
set.seed(162738)


# Load the necessary library
library(glmnet)

# Plot the lasso path with varying lambda values
plot(glmnet(gerxdumtrain,
             gercredtrain$V21,
             family="binomial", # binomial for logreg
             alpha=1),
      xvar = "lambda",
      label = TRUE)
```



```
# Perform cross-validation to select the optimal lambda
cvgerlasso = cv.glmnet(gerxdumtrain, gercredtrain$V21, family="binomial", alpha=1)

# Plot the cross-validation results
plot(cvgerlasso)
```



```
# Get the coefficients for the optimal lambda
coeflassoger = predict(cvgerlasso, new=gerxdumtest, s="lambda.min", type="coefficients")

# Display the coefficients and count non-zero coefficients
coeflassoger
```

```
## 49 x 1 sparse Matrix of class "dgCMatrix"
##          lambda.min
## (Intercept) -1.454739e+00
## V2          3.182095e-02
## V5          4.126319e-05
## V8          1.650686e-01
## V11         3.332994e-02
## V13        -6.391242e-03
## V16         3.849864e-02
## V19        -2.446448e-02
## V20         5.901246e-01
## V1_A12      -1.356402e-01
## V1_A13      -1.335883e+00
## V1_A14      -1.683169e+00
## V3_A31       4.151325e-01
## V3_A32      .
## V3_A33      .
## V3_A34      -5.694585e-01
## V4_A41      -6.424560e-01
## V4_A410     -5.624105e-01
## V4_A42      .
## V4_A43      -4.904551e-02
## V4_A44      .
```

```
## V4_A45      .
## V4_A46      1.322736e+00
## V4_A48      .
## V4_A49     -1.339864e-01
## V6_A62      .
## V6_A63      .
## V6_A64      .
## V6_A65     -8.433750e-01
## V7_A72      1.246955e-01
## V7_A73      .
## V7_A74     -1.695024e-01
## V7_A75      .
## V9_A92      .
## V9_A93      .
## V9_A94     -2.781214e-01
## V10_A102     .
## V10_A103    -6.721503e-01
## V12_A122     .
## V12_A123     .
## V12_A124     .
## V14_A142     .
## V14_A143    -3.515876e-01
## V15_A152    -9.237637e-02
## V15_A153     .
## V17_A172     .
## V17_A173     .
## V17_A174     .
## V18_2        .
```

```
length(coeflassoger[coeflassoger[,1] != 0 ,])
```

```
## [1] 26
```

We see that variable selection was performed.

```
# Make predictions on the test set using the selected lambda
predlassoger = predict(cvgerlasso, new=gerxdumtest, s="lambda.min", type="response")

# Display the first 10 predicted values
predlassoger[1:10]
```

```
## [1] 0.14378795 0.14396077 0.06977632 0.37842897 0.35945215 0.50497267
## [7] 0.46602492 0.43738247 0.31481853 0.01000166
```

- The lasso keeps 26 covariates (plus the intercept) out of the 48.
- The **predictions** are the **probabilities** of being a bad risk.
- Need a **threshold**.

We will use a function to estimate the threshold c which maximizes the **gain matrix**

$$G = \begin{pmatrix} g_{11} & g_{12} \\ g_{21} & g_{22} \end{pmatrix} = \begin{pmatrix} TP & FN \\ FP & TN \end{pmatrix}$$

$$\begin{aligned} \max \quad & \mathbb{P}(\hat{y} = 1, y = 1) \times g_{11} + \mathbb{P}(\hat{y} = 1, y = 0) \times g_{12} \\ & + \mathbb{P}(\hat{y} = 0, y = 1) \times g_{21} + \mathbb{P}(\hat{y} = 0, y = 0) \times g_{22} \end{aligned}$$

Cross-validated probabilities

First, we create a function to obtain cross-validated probabilities from `glmnet`. - The best λ is chosen by CV in each fold. - The output can be used to find the best threshold afterwards.

```
# Function to get cross-validated estimated probabilities from glmnet
# The best lambda is chosen by CV in each fold
# The output can be used to find the best threshold afterwards

predcvglmnet = function(xtrain, ytrain, k = 10, alpha = 1)
{
  # xtrain = matrix of predictors
  # ytrain = vector of target (0-1)
  # k = number of folds in CV
  # alpha = alpha parameter in glmnet

  # Load the necessary library
  library(glmnet)

  # Set a seed for reproducibility
  set.seed(375869)

  # Get the number of observations in the training data
  n = nrow(xtrain)

  # Initialize the vector to store predicted probabilities
  pred = rep(0, n)

  # Create a random permutation of indices
  per = sample(n, replace = FALSE)

  # Initialize indices for the current fold
  tl = 1

  # Perform k-fold cross-validation
  for (i in 1:k)
  {
    # Determine the upper index for the current fold
    tu = min(floor(tl + n / k - 1), n)

    # Adjust for the last fold
    if (i == k)
    {
      tu = n
    }

    # Get the current indices for the fold
    cind = per[tl:tu]

    # Fit a glmnet model with cross-validation on the current fold
    fit = cv.glmnet(xtrain[-cind, ], ytrain[-cind], family = "binomial", alpha = alpha)

    # Predict the probabilities for the current fold using lambda.min
    pred[cind] = predict(fit, new = xtrain[cind, ], s = "lambda.min", type = "response")

    # Update the starting index for the next fold
    tl = tu + 1
  }

  # Return the predicted probabilities
}
```



```

pred
}

```

Note that here there are two levels of cross validation:

1. The **outer CV** is used to compute **estimated probabilities**.
2. The **inner CV** estimates the tuning parameter for a given fold of the outer CV.

Best Binary Classifier Threshold

```

# Function to find the best threshold to use for a
# binary classifier with respect to a gain matrix

bestcutp = function(predcv, y, gainmat = diag(2), cutp = seq(0, 1, .02), plotit = FALSE)
{
  # predcv = vector of predicted probabilities (e.g., obtained out-of-sample by CV)
  # y = vector of target labels (0 or 1)
  # gainmat = gain matrix (2x2) (we want to maximize the gain)
  # (1,1) = gain if pred=0 and true=0
  # (1,2) = gain if pred=0 and true=1
  # (2,1) = gain if pred=1 and true=0
  # (2,2) = gain if pred=1 and true=1
  # cutp = vector of thresholds to try
  # plotit = whether to plot the results

  # Initialize variables
  nc = length(cutp)  # Number of thresholds to evaluate
  gain = rep(0, nc)  # Vector to store calculated gains

  # Loop through each threshold value
  for (i in 1:nc)
  {
    pred = as.numeric(predcv > cutp[i]) # Predicted binary outcomes using the threshold
    gain[i] = mean(gainmat[1, 1] * (pred == 0) * (y == 0) + # Calculate gain for this threshold
                  gainmat[1, 2] * (pred == 0) * (y == 1) +
                  gainmat[2, 1] * (pred == 1) * (y == 0) +
                  gainmat[2, 2] * (pred == 1) * (y == 1))
  }

  # Optionally plot the gains over different thresholds
  if (plotit)
  {
    plot(cutp, gain, type = "l", xlab = "Threshold", ylab = "Gain")
  }

  # Create a list containing results
  out = list(NULL, NULL)
  out[[1]] = cbind(cutp, gain) # Matrix of thresholds and associated gains
  out[[2]] = out[[1]][which.max(gain),] # Threshold with the maximum gain and its associated mean gain
  out
}

```

Finding the optimal cutoff using the gain matrix

```

# Computing CV estimated probabilities with glmnet
set.seed(16274)


```

```

pred = predcvglmnet(gerxdumtrain, gercredtrain$V21, k = 10, alpha = 1)

# Estimating the best threshold with the identity gain matrix
# This step is intended to find the threshold that maximizes gain
res = bestcutp(pred, gercredtrain$V21, gainmat = diag(2),
               cutp = seq(0, 1, .02), plotit = TRUE)

```



```

# Display the threshold with the associated mean gain
res[[2]]

```

```

##      cutp      gain
## 0.5400000 0.7716667

```

Computing the good classification rate

```

# using the best threshold to obtain the predictions
predlassoger01=as.numeric(predlassoger>res[[2]][1])

# good classification rate on the test set
mean(gercredtest$V21==predlassoger01)

```

```

## [1] 0.6975

```

```

# a naive rule would get a good classification rate of
max(mean(gercredtest$V21), 1-mean(gercredtest$V21))

```

```
## [1] 0.6775
```


We obtain a true good classification rate of 0.6975 on the test set, and a **naive rule**, assigning everyone to the 0 class would get a good classification rate of 0.677.

Hence, the lasso logistic regression performs only a little better compared to the naive rule.

AUC and ROC Curves

- The **ROC curve** is a plot of the **true positive rate** (TPR) against the **false positive rate** (FPR) for the different possible thresholds.
- The **AUC** is the area under the ROC curve.

```
# Load the ROCR library for ROC curve analysis  
library(ROCR)  
  
# Create a prediction object using predicted probabilities and true values  
predrocr = prediction(predlassoger, gercrtest$V21)  
  
# Calculate the ROC curve  
roc = performance(predrocr, "tpr", "fpr")  
  
# Plot the ROC curve  
plot(roc)  
  
# Add a diagonal reference line for a random classifier  
abline(a = 0, b = 1)
```



```
# Calculate and display the AUC (Area Under the Curve)
```

```
performance(predocr, "auc")@y.values[[1]]
```

```
## [1] 0.746074
```

- Note that we used the test data here but in practice, the value of Y would not be available for the new data.
- Hence, we would need to compute the ROC curve, AUC and lift chart with the training data.
- In that case, we must remember to use **proper estimation of the probabilities** (like the ones obtained by CV above), in order to get honest estimates.

lift curve

```
# Calculate and plot the lift chart
```


```
lift1 = performance(predocr, "tpr", "rpp")
```

```
# Plot the lift chart
```

```
plot(lift1)
```

```
# Add a diagonal reference line for a random classifier
```

```
abline(a = 0, b = 1)
```



Customizing the gain matrix


Instead of using the identity gain matrix, we can use a custom gain matrix to reflect the fact that the cost of a false positive is not the same as the cost of a false negative.

Ex.

$$G = \begin{pmatrix} 1 & 0 \\ 0 & 10 \end{pmatrix}$$

```
# Set the random seed for reproducibility
set.seed(18965)

# Estimate the best threshold with a gain matrix that favors detecting bad risks
res1 = bestcutp(pred, gercredtrain$V21, plotit = TRUE, gainmat = rbind(c(1,0), c(0,10)))
```



```
# Display the threshold with the associated mean gain
res1[[2]]
```

```
## cutp gain
## 0.120 2.935
```

In this case the optimal threshold is clearly much slower, since we now want to **classify more people as bad risks** because the reward is higher if we are right.

```
# Function to compute the C-index with a binary target
cindexbasic=function(phat,y)
{
  n=length(phat)
  cc=0
  npair=0
  for(i in 1:(n-1))
  {
```

```

for(j in (i+1):n)
{
  if(y[i]!=y[j])
  {
    cc=cc+(phat[i] > phat[j])*(y[i]>y[j])+(phat[i] < phat[j])*(y[i]<y[j])+ 0.5*(phat[i]==phat[j])
    npair=npair+1
  }
}
}
cc/npair
}

# We get the same value as the AUC
cindexbasic(predlassoger,gercredtest$V21)

```

```
## [1] 0.746074
```

And we see that it is indeed the same as the AUC.